


Regime detection that can say what happened — and a measurement of how much that is worth

An auditable explanation layer for any regime model. BSc Mathematics with Statistics, University of Bristol.

Finding: the explanations are specific and faithful to their sources; whether they are *diagnostic* of a regime change is unresolved, and this design could never have resolved it.



PROBLEM STATEMENT

A regime model tells you the market changed, not what happened — it only ever sees returns. So its output is coloured bands that a human interprets from memory. On a desk of thirty portfolio managers that is thirty private, unrecorded readings of one signal, and the reasoning leaves with whoever remembers the week.

SOLUTION OVERVIEW

Above. A two-state HMM marks *when* the regime changed; a language model reading *only news published before that date* describes what was happening. The two never mix: no statistic comes from the model, and it never moves a regime boundary.

The boundary is enforced by the type system — a window object refuses to construct if any item post-dates it. Pages are pinned to their revision id, because a mean of 38.2% of a Wikipedia day-page as it stands today (peak 71.2%) was written *after* the day it describes. Fetching the live page feeds the model hindsight, invisibly.

USE OF AI

claude-opus-5 via the Anthropic Messages API — one call per window, effort high, output schema enforced server-side. Prompts live in version-controlled files with a dated iteration history, never inline strings; every call logs the model id, prompt hash and input hash, so any sentence in the results traces to the exact prompt and news window behind it. Built with Claude Code. Wikipedia MediaWiki API, yfinance, hmmlearn, scikit-learn, Streamlit.

IMPACT & VALUE

It replaces “trust me, that band is the Greek referendum” with a written explanation citing dated sources — **arriving with its own error bars**. Tested on 20 transitions and 40 matched control dates:

Can you trust the explanations?	What we found
Date hidden: can we still tell which two weeks it describes?	72% — guessing gives 33%
Claims really come from the source cited	96.4% — a random source gives 1.8%
Invented citations	none, in 474 claims
Confident on real events, vs on ordinary days	80% vs 62% — ignoring the news gives no gap
Whole study re-run on a different AI	same scores

Why this is not a chatbot. Anyone can ask a model what happened in a month and get a fluent answer — and it reads identically whether it is specific to that month or would fit any month. **A chatbot gives you an answer; this gives you an answer and how far to trust it.** The reusable asset is the controls, not the explanations: point it at any regime model's dates, any asset, any method, and every explanation returns with that scorecard attached.

REFLECTIONS

Row four is the finding. The model explains an ordinary Tuesday almost as readily as a real transition. Calling that “not diagnostic” would be a no-difference claim $p = 0.43$ does not license: it is +18.5pp, 95% CI [-12, 47]pp, unresolved — power at that gap is only 0.17, and the constraint is 20 transitions, not controls, so more placebos could never have helped. **What this buys today is audibility, not signal generation.**

Next: pre-specify volatility *onsets* rather than all transitions, which roughly triples the power. One leak stays open: the boundary is 23:59 UTC and Wikipedia reports the US close the same evening.